

Project Documentation: Cloud Infrastructure for Recipe Storage & Distribution

Overview

This project provisions a secure and scalable infrastructure using AWS services for storing, archiving, processing, and distributing recipe files globally. The infrastructure is built with Terraform, ensuring automation, version control, and reproducibility.

AWS Cloud Services Used & Their Functionality

1) Amazon S3 (Simple Storage Service)

- **S3 Bucket for Recipe Storage** (automation-recipe-storage)
 - Stores uploaded recipe files.
 - Versioning enabled to keep track of file changes.
 - Server-side encryption (AES-256) to secure data.
 - **S3 Bucket for Archives** (automation-recipe-archives)
 - Stores processed and archived recipe files.
 - Versioning enabled for tracking changes.
- 2) S3 Lifecycle management for both buckets for effective cost optimization
- Files in S3 bucket **automation-recipe-storage move to Glacier** after 30 days.
 - Files in recipe_archives **expire after 1 year**.

3) AWS Lambda

Used for: Processing recipe files (e.g., zipping and moving them to archives)

- Lambda function (zip_recipes)
- Triggered when a new file is uploaded to recipe_storage.
- Compresses the file and stores it in recipe_archives.
- Uses an **IAM Role** (lambda_exec) with permissions to read/write S3.

Lambda Function Role in the Workflow

- A user uploads a new recipe file → recipe_storage S3 bucket.
- **S3 Event Notification** triggers **Lambda** automatically.
- **Lambda downloads the file**, compresses it, and **uploads the ZIP** to recipe_archives.
- **CloudFront** serves archived files securely to users.

- **S3 Event Notification**
- **S3** triggers **Lambda** when a new file is uploaded.
- **Uses S3 bucket** notification to invoke Lambda.

4) AWS IAM (Identity & Access Management)

Used for: Securing access to AWS resources

- **IAM Role for Lambda (lambda_exec)**
- Allows Lambda to interact with S3 and CloudWatch.
- IAM Policy for Lambda to access S3
- Grants Lambda permissions to:
- Read/write recipe files in both S3 buckets.

IAM Policy for Lambda to write logs

- Allows Lambda to log execution details in CloudWatch.
- S3 Bucket Policy
- Allows CloudFront to securely retrieve archived recipe files.

5) Amazon CloudFront

Used for: Secure global distribution of archived recipe files

- **CloudFront Distribution (recipe_distribution)**
- Provides fast, secure access to recipe archives.
- Uses **Origin Access Identity (OAI)** to restrict direct access to S3.
- **HTTPS enforced** for security.

6) AWS CloudWatch

Used for: Monitoring & Logging

- CloudWatch logs Lambda execution for debugging and monitoring in the log group `/aws/lambda/zip_recipes`

Key Features & Functional Flow of the Project

1. **File Upload:** Users upload recipe files to `recipe_storage` S3 bucket.
2. **File Processing:**
 - S3 triggers Lambda when a new file is uploaded.
 - Lambda compresses the file and moves it to `recipe_archives`.
3. **Lifecycle Management:**
 - Files in `recipe_storage` move to Glacier after 30 days.
 - Files in `recipe_archives` expire after 1 year.
4. **Secure Distribution via CloudFront:**
 - Clients access archived recipes through cloudfront.
 - OAI ensures direct S3 access is blocked.

Security Considerations

- IAM Policies restrict access to only necessary actions.
- S3 Encryption ensures data is protected at rest.
- CloudFront OAI prevents unauthorized S3 access.

Note > "In this project, I opted not to use CloudFront Signed URLs for Access Control and AWS Site-to-Site VPN. While I initially considered implementing Site-to-Site VPN, I found the configuration process to be time-consuming. As a result, I prioritized other AWS services and deferred its implementation through Terraform for the time being."

Additionally, we encourage you to provide brief insights into the benefits, potential pitfalls, and important considerations surrounding the following topics:

- 1) IaC
- 2) IoT and Edge computing
- 3) Cost calculations and optimisations

1. Infrastructure as Code (IaC)

- **Benefits:**

- Automation & Consistency – Automates infrastructure deployment, reducing human error.
- Scalability – Easily replicates infrastructure across environments (dev, test, production).
- Version Control – Treats infrastructure like code, enabling rollback and tracking changes.
- Faster Deployments – Speeds up infrastructure provisioning with templates and scripts.

- **Pitfalls:**

- Complexity – Learning IaC tools (Terraform, CloudFormation, Pulumi) can be challenging.
- Misconfigurations – A small mistake in code can deploy incorrect resources at scale.
- State Management Issues – Keeping track of infrastructure state across teams requires careful handling.

Important Considerations:

- Use GitOps to manage IaC changes efficiently.
- Implement automated testing (e.g., terraform plan, policy checks) before applying changes.
- Follow the principle of least privilege when managing IAM permissions in IaC.

2) Internet of Things (IoT) and Edge Computing

Benefits:

- Low Latency – Edge computing processes data closer to the source, reducing response times.
- Bandwidth Savings – Less data is sent to centralized cloud services, reducing costs.
- Reliability – Works even with intermittent connectivity by processing data locally.
- Scalability – Supports a vast network of IoT devices in real-time applications.

Pitfalls:

- Security Risks – More entry points mean increased attack surfaces for hackers.
- Device Management Complexity – Keeping firmware and software updated across devices is challenging.
- Data Privacy Concerns – Edge devices handling sensitive data may require regulatory compliance (GDPR, HIPAA).

Important Considerations:

- Use strong encryption (TLS, AES) for data transmission and storage.
- Implement zero-trust security models for device authentication.
- Optimize data filtering at the edge to reduce cloud processing and storage costs.

3) Cost Calculations and Optimizations

Benefits:

- Cost Transparency – Helps businesses understand where cloud spend is going.
- Better Budgeting – Enables accurate forecasting of cloud expenses.
- Optimized Resource Usage – Identifies underutilized resources and rightsizes them.
- Cost Savings – Helps reduce unnecessary expenses via reserved instances, auto-scaling, and spot instances.

Pitfalls:

- Unexpected Costs – Misconfigured auto-scaling or forgotten resources can lead to bill shock.
- Complex Pricing Models – Cloud providers have intricate pricing structures that are hard to track.
- Over-provisioning – Allocating more resources than needed leads to waste.

Important Considerations:

- Use AWS Cost Explorer, AWS Compute Optimizer, Azure Cost Management, or GCP Billing Reports to track spending.
- Implement auto-scaling to match demand and avoid over-provisioning.
- Leverage spot instances and reserved capacity for predictable workloads.
- Regularly review idle resources (unused volumes, old snapshots) and delete unnecessary ones.

